

Functional Requirements

The system has to count video view counts

WRITE

countViewEvent(videoId)

countView(videoId, eventType{VIEW|LIKE|SHARE})

processEvent(videoId, eventType, function{COUNT|SUM|AVERAGE})

processEvent(listOfEvents)

The system has to return video views count for a time period

READ

getViewCount(videoId, startTime, endTime)

getCount(videoId, eventType, startTime, endTime)

getStats(v)deold, eventType, function, startTime, endTime)

Non-Functional Requirements

Scalable (10K's request per second)

Highly Performant (few 10's of milliseconds to return view count of a video)

Highly Available (survives hardware, network failure, no single point of failure)

Design

Tables

TableName: Events

Coulumns: videoId, timestamp

Notes

- Fast Writes

- Can recelculate when needed

- Slow reads

- Costly when too many events

- Used with batch data processing

- Keep events for some time and purge events

Table Name:Count_per_minute

Columns: videoId, lastViewTimestamp, count

Notes

- Aggregated data

- Query the way data was aggregated

- Required a data aggregation pipeline

- Used with stream processing

Relational DB

TableName: Video_info

Columns: id, name, channel_id

TableName: video_stats

Columns: id, timestamp, count

TableName: Channel_info

Columns: id, name

Non-Relational DB

Video_id, channel_name, video_name, 15:00, 16:00, 17:00

Ingestion Path Components

- **API Gateway/PartitionerServiceClient**

Blocking vs Non-Blocking IO: Blocking systems are easy to debug

Buffering & Batching

- Put events into buffering for several seconds or until batch fills up
- Events are combined in API Gateway before sending to Partitioner Service
- Increases throughput
- Saves Cost
- Request Compression is more effective

Timeouts

Connection Timeout (10ms)

Request Timeout: Measure latency Percentile

Retries

All retries at same time may cause retry storm & overload server with too many requests. Use **Exponential Backoff & Jitter**

Circuit Breaker Pattern; Calculate failed requests count in given amount of last many minutes & stop calling downstream service. Then allow limited request for some time if succeeds allow all requests. Difficult to test, Hard to set error threshold and timeout

- **LoadBalancer**

Software vs Hardware Load balancers

Hardware load are devices that can be bought and handle millions of request

Network protocols

TCP: load balancers forward packet without inspecting contents of packets(Super fast can handle millions of request per second)

HTTP: terminates the connection, Can look into contents of message, can make load balancing decision based on contents of packets

Load Balancing Algorithms

- Round Robin - Inorder across lists of servers
- Least Connection
- Least Response Time

- Hashed Based: On Client IP or URL

DNS: Register Partitioner Service in DNS

Health Checking: Are servers healthy

High Availability: Primary & Secondary load balancers maybe in different datacenters

- **Partitioner Service & Partitions**

Partition Strategies:

- 1) Hasing: Can cause hot partition for large scale systems
- 2) Consistent Hashing: Dedicated Partitions

Service Discovery (Zookeeper): getPartitionsListByPartitionService()

Single leader Replication

Message Formats

Text

- 1) XML
- 2) CSV
- 3) JSON

Binary (For large scale rel time systems)

- 4) Thrift
- 5) Protocol Buffer
- 6) Avro
- 7) Parquet

Faster to parse, Schema can be shared in DB from where producer & consumer can retrieve them

Technology Stack

API Gateway/PartitionServiceClient

Netty

Netflix hystrix

Load Blancing

NetScaler(Hardware LB)

NGINX

AWS ELB

Messaging System

Apache Kafka

AWS Kinesis

Data Procssing

Apache spark (Stream processing)

Apache Flink

AWS Kinesis Data Analytics

Storage

Apache Cassandra

Apache HBASE

InfluxDB

AWS S3 (DFS for Rollup data)

Caching

- Redis

- AWS ElasticCache

DeadLetters Queue

- AWS SQS

- RabbitMQ

Embedded DB

- RocksDB

Leader Election

- Zookeeper

Service Discovery

- Zookeeper

- Netflix Eureka

Monitoring

- AWS Cloudwatch

- ELK

Message Format

- AVRO

Partition Data Algorithm

- Murmur Hash

Testing

- Apache Jmeter

Testing

- 1) Performance testing: under heavy load to identify bottlenecks
- 2) Load testing: Behavior under specific expected load
- 3) Stress Testing: Beyond operational capacity often to breaking point. What will break first
- 4) Leaks in resources: Typical production load for expected period of time

Health Monitoring

- Metric, Dashboard, Alarms, Latency, Traffic, Errors, Saturation