

## Functional Requirements

sendMessage(messageBody)

receiveMessage()

## Non-Functional Requirements

Scalable (Handle load with increasing queues and messages)

Highly Available (Survives hardware & network failure)

Highly Performant (single digit latency for operations)

Durable (data persisted when submitted to queue until polled)

SLA (minimum throughput)

## Design

### Design Patterns

- Bulkhead Pattern
- Circuit Breaker Pattern

### Components

- **FrontEnd WebService**
  - **Lightweight web service**
  - **Stateless service deployed across data centers**
  - **Actions**
    - Request validation
      - Required params are present (e.g. Queue name)
      - Message Size not exceeding threshold value
    - Authentication / Authorization
    - TLS/SSL termination
      - SSL termination on load balancer is expensive
      - Usually handled by a separate TLS HTTP proxy runs as a process on same host
    - Caching (Server Side)
      - Cache stores copies of source data
      - Stores info about most frequently used queue to save cost on authentication and authorization service
      - It helps to reduce load to Backend services, increases overall system throughput & availability, decreases latency
      - Stores previously seen requestIds
    - Rate Limiting (Throttling)
    - Request Dispatching
      - Responsible for all activities associated with sending request to Backend service like client management, response handling, resource isolation etc.
    - User data collection
    - Server Side encryption
      - Messages are encrypted as soon as received

- Messages are stored in encrypted format and FrontEnd decrypt message only when they are sent back to client
- Request Deduplication
  - May occur when a response from a successful sendMessage(messageBody) failed to reach client
- Delivery Semantics
  - Lesser an issue for 'at least once', a bigger issue for 'exactly once' & 'at most once'
  - Caching is usually done to store previously seen requestIds
- **Metadata Service**
  - Stores information about queues
  - Many reads, little writes
  - Strong consistency storage is preferred but not required
  - Sharded to handle load
- **BackEnd Service**
  - Data stored in RAM and disk of backend host
  - FrontEnd Service retrieves BackEnd Service host information from Metadata Service
  - Replicated with in groups of host
    - a) Leader-Follower relationship
      - Incluster manager (QueueId, Leader Host, Followers)
        - Manages queue assignment within cluster
        - Maintain a list of host in the cluster
        - Monitors heartbeat from hosts
        - Deals with leader & follower failures
        - Splits queue between cluster nodes
    - b) Small clusters of independent hosts
      - Outcluster Manager (QueueId, ClusterId)
        - Manages queue assignment among clusters
        - Maintain a list of clusters
        - Monitors each cluster health
        - Deals with overrated cluster
        - Splits queue between clusters

Queue creation & deletion

Message deletion and mark invisible like SQS

Message Replication Synchronous & Asynchronous

Message delivery semantics atleast once is hard

Push vs Pull vs Long Polling

FIFO

Security (SSL over HTTP, encrypt in backend service)

Monitoring - FE, BE, Metadata Service, Log Data, Alerts/Alarms, Customer Dashboard